

Automaton Auditor — Final Report (Saturday Submission)

Repository: <https://github.com/78gk/The-Automaton-Auditor>

Date: 2026-02-28 (Saturday — Final Submission)

Self-audit overall score: **3.7 / 5.0**


Executive Summary

The Automaton Auditor is a hierarchical LangGraph StateGraph swarm that performs forensic auditing of AI-generated codebases using a Digital Courtroom design: Detectives collect typed evidence, Judges debate under conflicting personas, and a deterministic Chief Justice synthesizes the final verdict. This final report summarizes the Saturday self-audit results and provides an actionable remediation plan to maximize rubric alignment.

Architecture Deep Dive (Fan-Out / Fan-In + Dialectics)

Data Flow: ContextBuilder → [Detectives in parallel] → EvidenceAggregator → [Judges in parallel] → ChiefJustice → Markdown report. Dialectical synthesis is implemented via Prosecutor/Defense/TechLead personas and resolved using deterministic governance rules (security override, fact supremacy, functionality weight, variance dissent).

Architecture Diagram (StateGraph)



Self-Audit Criterion Breakdown

Final Chief Justice scores per rubric dimension (see full judge arguments in `audit/report_onself_generated/audit_report.md`):

Dimension	Final Score (/5)
Git Forensic Analysis	4
State Management Rigor	4
Graph Orchestration Architecture	4
Safe Tool Engineering	4
Structured Output Enforcement	4
Judicial Nuance and Dialectics	3
Chief Justice Synthesis Engine	4
Theoretical Depth (Documentation)	4
Report Accuracy (Cross-Reference)	3
Architectural Diagram Analysis	3

Prioritized Remediation Plan

- 1) Report Accuracy:** Update the final PDF to reference only real repo paths and include an explicit verified vs hallucinated table.
- 2) Swarm Visual:** Update the diagram to label all nodes and clearly show both fan-outs and fan-ins.
- 3) Judicial Nuance:** Further enforce persona divergence constraints in judge prompts.